AI-Based Stroke Phase Detection and Symmetry Analysis in Competitive Swimming Footage: A Proof of Concept

Daniel Pace

Supervisor: Dr. Vanessa Camilleri

29th April 2026

Submitted in partial fulfilment of the requirements for the degree of Bachelor of Science in Information Technology (Honours) (Artificial Intelligence).

L-Università ta' Malta
Faculty of Information & Communication Technology

Abstract

Computer vision analysis in aquatic environments presents unique challenges due to refraction, turbulence, and occlusion, which hinder the accuracy of standard motion tracking systems. This project proposes a non-intrusive AI-based proof of concept system for reliable pixel-to-meter scaling using markers in the background. The core contribution is a novel methodology for dynamic pixel-to-meter scaling using standard lane rope markers as calibration objects, addressing the limitations of fixed-camera setups in varying pool environments. This progress report outlines the development of the modular marker prototypes and the architectural design of the computer vision pipeline. It concludes by defining the evaluation protocols and the remaining milestones required to validate the system's robustness in real-world scenarios

Acknowledgements

Thanks

Contents

Abstract **i**

Acknowledgements **ii**

Contents **iv**

List of Figures **v**

List of Tables **vi**

List of Abbreviations **1**

Glossary of Symbols **1**

1 Introduction **1**
1.1 Problem Statement 1
1.2 Motivation 1
1.3 Aims and Objectives 1
1.4 Project Structure 2

2 Background **3**
2.1 Object Detection and Classification 3
2.2 Semantic Segmentation 3
2.3 Human Pose Estimation 4
2.4 Motion Compensation and Ego-motion 5
2.5 State Estimation and Trajectory Smoothing 5
2.6 Summary 5

3 Literature Review **6**
3.1 Evolution of AI in Sports Analytics 6
3.2 An Adversarial Domain for Computer Vision 6
3.3 From Handcrafted Features to Foundation Models 6
3.4 The Challenge of Metric Scaling and Calibration 7
3.5 Addressing the Data Scarcity Gap: Model-in-the-Loop Annotation 7

3.6 Summary of Research Needs 8

4 Methodology **9**
4.1 Marker Design and Visibility 9
4.2 Dataset Generation and Model-in-the-Loop Annotation 10
4.2.1 Automated Labelling with SAM 2 10
4.2.2 Outlier Filtering 10
4.2.3 Dataset Selection 11
4.2.4 Human-in-the-Loop Verification 11
4.3 Dynamic Scale Extraction Pipeline 11
4.3.1 Marker Detection via YOLO26 12
4.3.2 Refraction-Aware Scaling 13
4.4 Kinematic Metric Extraction 13
4.4.1 Swimmer Tracking via ViTPose 13
4.4.2 Velocity and Kinematic Profile Extraction 13
4.5 Summary 14

5 Evaluation **15**
5.1 Evaluation Plan 15
5.1.1 Quantitative Evaluation 15
5.1.2 Qualitative Validation 15
5.2 Evaluation Plan 15
5.2.1 Evaluation Plan 15

6 Conclusion **17**
6.1 Conclusions including Expected Outcomes and Difficulties/Challenges . . 17
6.1.1 Aims and Objectives 17

A Sample A **20**

B Sample B **21**

List of Figures

Figure 4.1 An image of the final marker design in use. 9
Figure 4.2 Marker Detection Model Training Pipeline. 12

List of Tables

1 Introduction

1.1 Problem Statement

Performance analysis in competitive swimming is traditionally reliant on manual observation or expensive, sensor-based systems that can be intrusive to the athlete [1] While computer vision offers a non-intrusive alternative, the aquatic environment introduces significant noise. Light refraction at the air-water interface, surface turbulence caused by the swimmer, and occlusions like bubbles and splashes degrade the quality of visual data [2, 3]. Furthermore, environment conditions, camera distances and angles vary significantly between venues, making consistent metric quantification, such as speed in meters/second, difficult without distinct calibration markers.

1.2 Motivation

There is a distinct lack of standardized, automated systems for extracting swimming metrics from video footage that do not require complex setups. A standard, simple, and well documented system capable of extracting metric data from video with minimal setup could act as a proof of concept, to prove the viability of such systems in aquatic environments. This project focuses on validating feasibility and robustness rather than achieving state-of-the-art performance.

1.3 Aims and Objectives

The primary aim of this project is to develop a Proof of Concept system that enables the extraction of swimming metrics from video footage. The specific objectives are:

1. To design, prototype, and create a modular, under-water marker design that maximizes underwater visibility.
2. To create and annotate a small dataset of swimming footage using the designed markers in the background.
3. To establish a pipeline for dynamic pixel-to-meter scale extraction using the markers for calibration.
4. To utilize the calibrated scale to provide useful kinematic metrics such as swimmer velocity and trajectory.

1.4 Project Structure

The remainder of this report is structured as follows.

Chapter 2– Background This chapter provides an overview of the core concepts and technological advancements relevant to the project.

Chapter 3– Literature Review This chapter reviews existing research and methodologies in the field of computer vision for sports analytics, with a focus on swimming performance analysis. It identifies gaps in current approaches and situates the proposed system within the broader context of the use of AI in sports technology.

Chapter 4– Methodology This chapter details the proposed methodology for developing the proof-of-concept system, including the design of underwater markers, dataset creation, and the calibration pipeline for dynamic scale extraction.

Chapter 5– Evaluation This chapter presents the evaluation of the proposed system, including assessments of each step of the pipeline, and an overall assessment of the system's performance in extracting swimming metrics from video footage.

Chapter 6– Conclusion and Future Work This chapter summarizes the findings of the project, discusses the implications of the results, and outlines potential directions for future research and development in this area.

2 Background

Existing literature in sports biomechanics highlights the importance of kinematic metrics like acceleration and hand trajectory. While computer vision research has successfully applied deep learning for detection, classification, segmentation, and human pose estimation (HPE), a critical limitation in many CV-based swimming systems is the conversion of distances from the pixel space to real-world units, often relying on relative metrics or requiring external sensors.

This chapter covers the relevant foundations upon which the proposed system is built, including object detection and classification, semantic segmentation, and human pose estimation. The chapter concludes with a brief summary of the discussed techniques.

2.1 Object Detection and Classification

Object detection is a fundamental computer vision task that involves identifying the presence and location of objects within an image. In the context of swimming analysis, this typically involves localizing the swimmer and specific environmental markers. Modern detectors have shifted from region-based convolutional neural networks (R-CNNs) to "one-stage" detectors to optimize for speed and real-time performance.

The YOLO (You Only Look Once) series represents the state-of-the-art in one-stage detection. The iterations used in this work, YOLOv8 and YOLO26 [4], introduces an anchor-free detection mechanism. Unlike previous versions that relied on predefined anchor boxes, anchor-free models predict the centre of an object directly, which reduces the complexity of the non-maximum suppression (NMS) stage. This results in improved performance on varying scales, which is a critical feature when tracking markers that change size due to camera distance or water refraction. In sports analytics, YOLO-based architectures are favoured for their ability to process high-frame-rate video without significant latency [5].

2.2 Semantic Segmentation

Semantic segmentation involves assigning a categorical label to every pixel in an image, facilitating a more granular understanding of the scene than bounding boxes. For aquatic environments, segmentation is used to isolate the swimmer's body from highly textured backgrounds or to precisely delineate markers amidst bubbles and surface turbulence.

A significant advancement in this field is the Segment Anything Model (SAM) [6]. SAM is a "foundation model" trained on over a billion masks, enabling zero-shot gener-

alization to new domains, such as underwater footage, without additional training. This project utilizes SAM 2 [7], which extends this capability to the temporal domain through a memory-based attention mechanism. SAM 2 treats video as a continuous stream, using a "memory bank" to store features from previous frames. This allows the system to maintain a stable segmentation of markers even during transient occlusions or periods of significant water splash, addressing a core problem identified in the aquatic environment.

In addition to its role in tracking, SAM 2 was utilized in this research as a semi-automated annotation tool to facilitate dataset creation. Because manual pixel-level annotation of video sequences is prohibitively time-consuming, SAM 2's temporal propagation feature was used to generate high-fidelity masks for underwater markers across several video sequences from a few manual prompts. These generated masks were then converted into bounding boxes and used to fine-tune a custom YOLO26 model for marker detection. This "model-in-the-loop" approach enables the creation of robust, domain-specific detectors with significantly reduced manual labelling effort. It bridges the gap between general-purpose foundation models and specialized sports analytics tasks.

2.3 Human Pose Estimation

Human Pose Estimation (HPE) is the process of identifying and tracking the coordinates of specific anatomical keypoints, such as joints and extremities. In swimming, HPE is used to calculate stroke frequency and joint angles. Traditionally, HPE utilized High-Resolution Nets (HRNet) to maintain high-resolution representations throughout the network.

However, recent research has transitioned toward Transformer-based architectures. ViTPose [8] utilizes a non-hierarchical Vision Transformer (ViT) [9] as its backbone. Unlike convolutional neural networks (CNNs) that process images through local receptive fields, Transformers use self-attention to capture long-range dependencies across the entire frame. This is particularly advantageous in swimming, where a swimmer's limb may be extended far from the torso; ViTPose can leverage the global context of the swimmer's body to accurately predict the position of a wrist or ankle that might be partially obscured by water refraction. This project utilizes the large-scale variant of ViTPose to ensure high precision in keypoint localization before converting these coordinates into metric data through marker-based linear scaling and baseline projection.

2.4 Motion Compensation and Ego-motion

In sports videography, unless the camera is rigidly fixed to a stationary mount, the resulting footage contains "ego-motion"—displacement caused by the movement of the camera itself. For swimming analysis, where a coach may follow the athlete along the poolside, this introduces a significant challenge: the pixel coordinates of the swimmer are a summation of the swimmer's actual velocity and the camera's displacement.

To resolve this, computer vision systems utilize stationary environmental features, such as lane ropes or pool-floor markings, as a global coordinate system. By tracking these fixed points using a homography matrix or simpler translation vectors [10], the system can perform background registration. This process effectively "stabilizes" the coordinate space, allowing for the extraction of the swimmer's true kinematic trajectory relative to the pool rather than the image frame.

2.5 State Estimation and Trajectory Smoothing

Raw data extracted from deep learning models often contains high-frequency noise or "jitter" caused by partial occlusions, light refraction, or fluctuating confidence scores in the model's output. To transform these noisy observations into a physically plausible trajectory, state estimation techniques are employed.

The Kalman Filter [11] is the standard recursive algorithm for this purpose. It operates in two phases: a "predict" step, which estimates the current state (position and velocity) based on previous data and Newtonian laws of motion, and an "update" step, which adjusts this prediction based on the new noisy measurement from the CV model. This approach is particularly effective in swimming because it can bridge gaps in detection—such as when a swimmer's hand is completely submerged or hidden by a splash—by relying on the predicted momentum of the limb until it reappears.

2.6 Summary

The integration of YOLO for rapid detection, SAM 2 for robust temporal segmentation and automated dataset generation, and ViTPose for high-fidelity pose estimation forms a comprehensive, modular pipeline for swimming analysis. These tools collectively address the core challenges of the aquatic environment—namely visual noise, light refraction, and the historical lack of labeled underwater data—by leveraging recent breakthroughs in anchor-free detection, memory-based attention, and global spatial transformers.

3 Literature Review

3.1 Evolution of AI in Sports Analytics

Computer vision (CV) has fundamentally altered the landscape of sports analytics, transitioning from manual notation to automated, high-fidelity measurement. While early systems relied on background subtraction and handcrafted features, the current paradigm is dominated by Deep Learning (DL). As noted by [12], AI-driven systems are now the standard for extracting kinematic data and estimating pose in real-time.

However, a significant shift is currently underway, from task-specific CNNs to "Foundation Models", large-scale models trained on massive datasets that exhibit zero-shot generalization. This thesis utilizes models such as SAM 2 and ViTPose, which provide a level of geometric and spatial reasoning previously unavailable in sports-specific contexts.

3.2 An Adversarial Domain for Computer Vision

While field sports provide stable visual environments, the aquatic setting serves as an adversarial domain that stress-tests the limits of image processing. The primary challenge remains light refraction; the air-water interface creates non-linear geometric distortions that violate the fundamental pinhole camera model assumptions [13].

Beyond refraction, the environment is plagued by dynamic noise. As [3] highlights, measurements are strongly influenced by disruptive elements such as bubbles, splashes, and surface glare. These elements introduce severe, transient occlusions. Traditional trackers often fail in these high-turbulence zones. This necessitates the use of temporal memory mechanisms, such as the memory bank implemented in SAM 2 [7], which allows a system to remember the features of an object even when it is obscured by a splash or submerged beyond clear visibility.

3.3 From Handcrafted Features to Foundation Models

Historically, structural feature detection in swimming relied on the Hough Transform or colour-keying to isolate lane ropes or markers. However, these break down under varying pool illumination and water turbulence [14].

The subsequent shift toward Convolutional Neural Networks (CNNs), such as U-Net or earlier YOLO versions, offered better robustness to noise. For instance, [15] demonstrated the utility of YOLOv8 for athlete identification. Yet, CNNs are often lim-

ited by their local receptive fields, which can struggle with the long-range spatial dependencies required to identify a swimmer's full limb extension in a complex stroke.

This has led to the adoption of Vision Transformers (ViT). Models like ViTPose [8] utilize global self-attention mechanisms to process the entire frame simultaneously. This allows the model to leverage the context of the swimmer's torso to accurately predict the position of a wrist or ankle, even if that specific joint is visually distorted by surface ripples—a core technical choice reflected in this project's pipeline.

3.4 The Challenge of Metric Scaling and Calibration

A fundamental limitation in many CV-based swimming systems is the conversion of distance from pixel space to real-world, metric units. A velocity of x pixels per second is meaningless without a reference scale.

The gold standard remains 3D camera calibration, but the practicalities of poolside data collection often preclude this. Consequently, researchers have turned to planar homography using known environmental features. While [16] utilized lane ropes, recent work suggests that discrete, underwater markers provide a more stable reference. By establishing a fixed pixel-to-meter ratio based on known marker distances such as the 2.5m proposed in this thesis, systems can achieve metric accuracy without expensive sensor arrays. This project extends this logic by using SAM 2 to track these markers with high temporal consistency, ensuring the calibration scale remains valid across the entire video sequence.

3.5 Addressing the Data Scarcity Gap: Model-in-the-Loop Annotation

A persistent gap identified in the literature is the scarcity of publicly available, annotated underwater datasets. SwimmerNET [3] provides 2D keypoints but lacks underwater marker data. This creates a cold start problem for research.

A novel solution emerging in recent literature is Model-in-the-Loop (MITL) annotation. Rather than manual labelling, researchers use Foundation Models to "pseudo-label" datasets. By using SAM 2's zero-shot segmentation to identify markers in a few frames and propagating those masks through time, a high-quality training set can be generated for a more efficient detector like YOLOv8. This workflow, implemented in the current thesis, represents a departure from traditional manual annotation and addresses the research need for domain-specific underwater markers.

3.6 Summary of Research Needs

The reviewed literature confirms the following three persistent challenges. Traditional trackers lose the target during occlusions such as splashes, refraction renders simple linear scaling inaccurate, and that there is a lack of labelled data for underwater calibration features.

This research bridges these gaps by synthesizing SAM 2 for stable temporal segmentation, ViTPose for context-aware human pose estimation, and a custom YOLO26 detector trained via an automated MITL pipeline. By doing so, the project moves beyond qualitative video analysis toward a robust, metric-accurate proof-of-concept system for swimming biomechanical analysis.

4 Methodology

This chapter details the technical implementation of the proof-of-concept system. The methodology is structured around the four primary project objectives, transitioning from physical marker design to the final kinematic extraction engine.

4.1 Marker Design and Visibility

The first objective required a modular marker design capable of withstanding the high-turbulence environment of a swimming pool while remaining identifiable under varying refractive conditions.

The lack of existing documentation on underwater marker design for computer vision applications necessitated an iterative prototyping approach. This led to the decision to create a modular design for a 3d-printed marker that could be easily set up along the lane. The final design, as shown in Figure 4.1, consists of high-contrast, alternating black and white segments, which seamlessly screw into one another. This decision to use a modular design allowed us to test various configurations, mainly how many segments to stack, effectively altering the height of the marker.

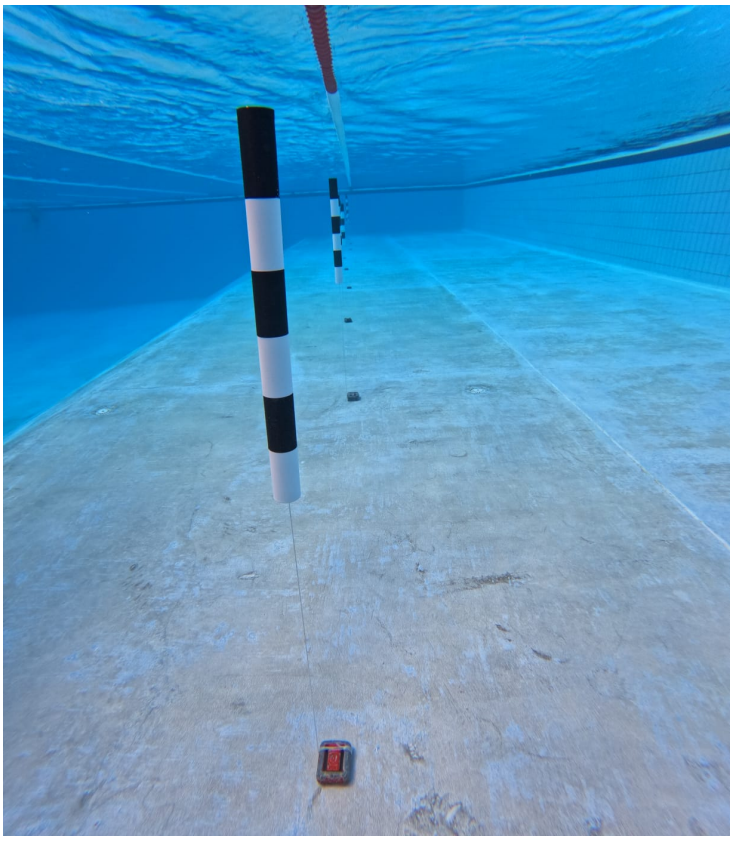

Figure 4.1 An image of the final marker design in use.

Each segment is 10 cm in height, with a 5 cm diameter, and the alternating pattern provides a unique visual signature that can be reliably detected by computer vision algorithms. The design also incorporates a screw mechanism that allows for easy assembly and disassembly, enabling quick setup along the lane for calibration purposes. This modularity not only facilitates testing different configurations but also ensures that the markers can be easily transported and stored when not in use.

The spacing between the markers was determined based on multiple tests, which revealed that an interval of 2.5m provided a good balance between visibility and the number of markers required for accurate calibration. This spacing allows for sufficient coverage along the lane while minimizing the number of markers needed, which is crucial for minimizing setup time and complexity.

4.2 Dataset Generation and Model-in-the-Loop Annotation

To address the second objective, necessitated by the lack of annotated underwater datasets, a "Model-in-the-Loop" (MITL) workflow was developed. This phase utilized foundation models to automate the labelling process for the downstream YOLO26 detector. The entire dataset creation pipeline is visualized in Figure 4.2.

4.2.1 Automated Labelling with SAM 2

The Segment Anything Model 2 (SAM 2) [7] was utilized to propagate marker bounding boxes across video sequences. Prompts were provided from the initial frames, propagated through the video. Once all tracked markers exited the frame, the process was repeated for the next chunk of frames until the entire video was annotated. This approach significantly reduced manual annotation time while maintaining high-quality labels. This process was repeated for multiple videos, resulting in a dataset of over 4,000 annotated frames, and over 7,000 bounding boxes.

4.2.2 Outlier Filtering

Automated tracking via foundation models, while efficient, is susceptible to "mask jitter" and drift, particularly when visual features are distorted by water refraction or transient occlusions. To ensure the robustness of the downstream YOLO detector, a statistical filtering stage was implemented to prune the raw SAM 2 output.

The filtering process focuses on the geometric consistency of the generated bounding boxes. Based on the assumption that the physical dimensions of the markers are constant, any significant variation in the predicted box area or aspect ratio was flagged as tracking noise. The system utilizes the Interquartile Range (IQR) method [17] to identify

and remove these outliers. For a set of bounding box areas A, the filtering bounds are defined as:

$$[\text{Lower Bound}, \text{Upper Bound}] = [Q_1 - 1.5 \times \text{IQR}, Q_3 + 1.5 \times \text{IQR}] \tag{4.1}$$

where Q_1 and Q_3 represent the first and third quartiles, respectively. Any frame containing boxes falling outside these bounds, often caused by SAM 2 "bleeding" into a splash or expanding to cover the background, was excluded from the training set. This ensured that the model was trained only on high-confidence, geometrically plausible representations of the underwater markers.

4.2.3 Dataset Selection

To ensure the large YOLO26 model was trained using the best possible dataset, an A/B testing approach was implemented. Two datasets were created: Dataset A, which included all SAM-generated labels, and Dataset B, which included only the filtered labels after the IQR outlier removal process. Both datasets were used to train separate YOLO26n models, and their performance was compared using standard metrics such as mean Average Precision (mAP) and visual inspection of the predicted bounding boxes. The dataset that yielded the best performance was selected for the final training of the final YOLO26s model.

4.2.4 Human-in-the-Loop Verification

To ensure the high precision required for metric extraction, a custom-developed graphical verification tool was utilized to perform a final audit of the filtered dataset. This tool, implemented using the Tkinter framework, allowed for the rapid visual inspection of the SAM-generated labels.

The interface provided the functionality to manually adjust bounding box boundaries or delete erroneous detections that bypassed the statistical IQR filter. This hybrid approach of combining foundation model speed with human precision addressed the "domain-shift" problem inherent in aquatic footage, where surface glare can occasionally mirror markers. This step was crucial for establishing a ground truth used for the final model fine-tuning and the subsequent accuracy evaluations in Chapter 5.

4.3 Dynamic Scale Extraction Pipeline

Objective 3 focuses on establishing a stable reference for converting image coordinates into metric units.

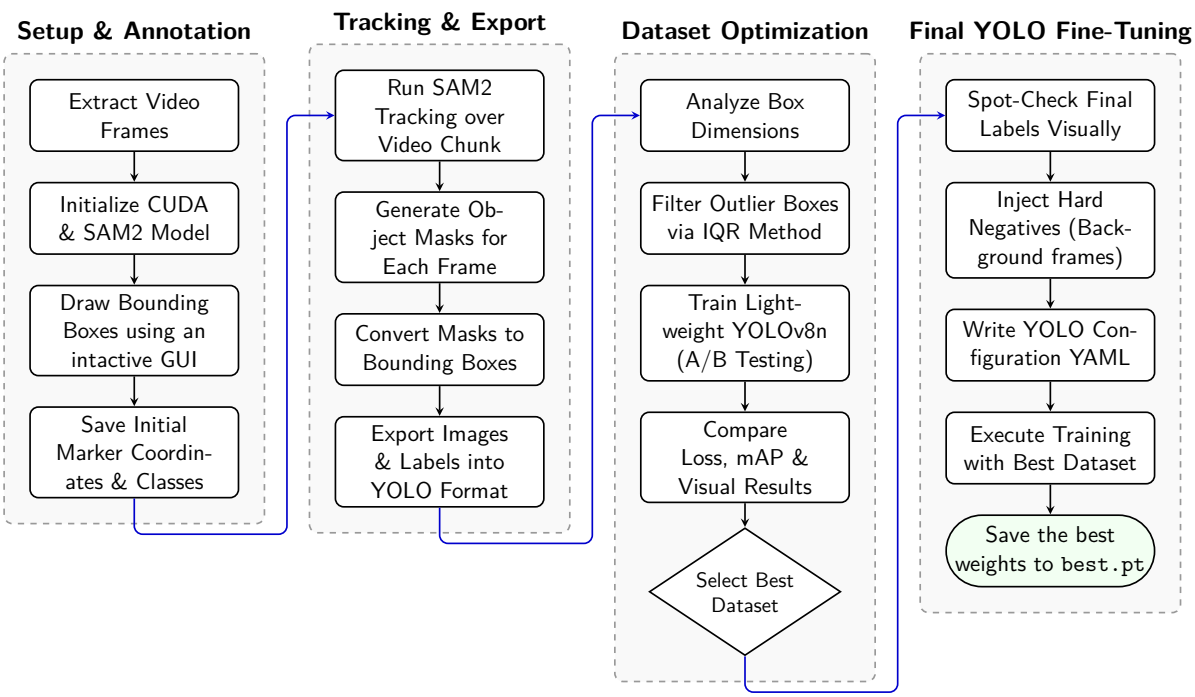

Figure 4.2 Marker Detection Model Training Pipeline.

4.3.1 Marker Detection via YOLO26

A custom YOLO26s model [4] was fine-tuned on the SAM-generated dataset to provide a robust method for localizing markers. To ensure high localization accuracy, which is critical for the subsequent pixel-to-meter ratio calculation, the training process utilized a Distribution Focal Loss (DFL) gain of 1.5 and a Box Loss gain of 7.5. These settings prioritize bounding box regression over classification accuracy, as the project involves a single-class detection task.

Training was conducted over 200 epochs with an image resolution of 640×640 pixels. To mitigate overfitting on the SAM-generated dataset, several regularization and augmentation strategies were employed:

- **Optimization Strategy:** An initial learning rate of 0.01 was used with a 3-epoch warm-up period and a momentum of 0.937 to stabilize early gradient descent.
- **Spatial Augmentations:** The model utilized Mosaic augmentation (1.0) and Random Erasing (0.4) to simulate occlusions and varying marker densities, forcing the network to learn robust features rather than memorizing specific marker placements.
- **Colour Distortions:** Variations in lighting and glare were accounted for by adjusting HSV (Hue, Saturation, Value) gains ($hsv_h = 0.015, hsv_s = 0.7, hsv_v = 0.4$), ensuring the detector remains effective under inconsistent illumination.

- **Early Stopping:** A patience of 20 epochs was implemented to terminate training if the validation fitness failed to improve, ensuring the model generalizes well to real-world frames.

The resulting model outputs high-confidence bounding boxes that serve as the foundation for physical distance estimation, leveraging the known dimensions of the physical markers to establish spatial scale.

4.3.2 Refraction-Aware Scaling

As identified in preliminary experiments (Section 4.2), frame-independent scaling is susceptible to refractive noise. To stabilize the scale, the system implements a Median Absolute Deviation (MAD) inlier mask [18]. This filter calculates the pixel-to-meter ratio (px/m) for all visible marker pairs and rejects outliers that deviate significantly from the median, ensuring a stable calibration constant C for each frame.

4.4 Kinematic Metric Extraction

The final objective involves the synthesis of pose estimation and scaling logic to calculate performance metrics.

4.4.1 Swimmer Tracking via ViTPose

The swimmer's position is tracked using [8]. The system extracts the coordinates of the hip keypoints to estimate the centre of mass (CoM). This approach is more robust than simple bounding box tracking, as it remains focused on the athlete's core during the stroke cycle.

4.4.2 Velocity and Kinematic Profile Extraction

The instantaneous velocity (v) is calculated by taking the first derivative of the smoothed displacement (d) over time (t). To transform pixel displacement (Δp) into meters, the calibration constant C from Phase 3 is applied:

$$v = \frac{\Delta p \cdot C}{\Delta t} \tag{4.2}$$

To mitigate high-frequency jitter and observation noise, the raw tracking records undergo an initial robust outlier removal using a Median Absolute Deviation (MAD) mask and a temporal jump filter. The cleaned trajectory is then smoothed using a discrete-time Kalman Filter [11], which simultaneously estimates the swimmer's continuous position

and instantaneous velocity. From these filtered state estimates, kinematic profiles such as absolute speed, peak velocity, and mean velocity are derived for analysis.

4.5 Summary

The methodology follows a logical progression from the physical environment to digital metrics. By aligning each phase with the project objectives, the system provides a robust framework that accounts for the unique adversarial factors of the aquatic domain.

5 Evaluation

5.1 Evaluation Plan

The evaluation strategy is two-fold, encompassing a rigorous quantitative assessment of the core technology and a qualitative validation of the final derived metrics.

5.1.1 Quantitative Evaluation

- **Lanerope Accuracy Comparison**: The two lanerope pipelines (Traditional CV and CNN) will be assessed against annotated ground truth. Metrics will include axis error (deviation from the true lanerope angle), scale error (difference from the known float interval distance), and frame-to-frame jitter (stability).
- **Downstream Metric Accuracy**: The lanerope method yielding the lowest error will be used to calibrate the system. The final system's metric outputs (e.g., a calculated 5-meter distance) will be compared to a physically measured distance to quantify the final distance estimation accuracy.

5.1.2 Qualitative Validation

- **Annotated Footage Comparison**: System outputs (keypoints, trajectories, and phase segmentations) will be overlaid onto the original video and validated against human-annotated samples.
- **Expert Feedback**: Structured interviews and feedback sessions will be conducted with certified swimming coaches to assess the interpretability, practical value, and actionability of the generated performance reports (e.g., symmetry score, wall contact duration).

5.2 Evaluation Plan

The evaluation strategy is two-fold, encompassing a rigorous quantitative assessment of the core technology and a qualitative validation of the final derived metrics.

5.2.1 Evaluation Plan

To evaluate the performance of step 1, the detection rate, as well as Intersection over Union (IoU), will be calculated against a manually annotated ground truth dataset. For

step 2, the accuracy of the pose estimation will be assessed using standard metrics such as Percentage of Correct Keypoints (PCK) and Mean Per Joint Position Error (MPJPE) against annotated keypoints.

To evaluate step 3, the calculated metrics (velocity, stroke length) will be compared against ground truth measurements obtained from manual timing and distance measurements during controlled swimming trials.

6 Conclusion

6.1 Conclusions including Expected Outcomes and Difficulties/Challenges

This project is expected to deliver a robust proof-of-concept system capable of providing high-value, objective biomechanical data from standard video. The expected outcomes include the validated optimal lanerope calibration method (Traditional CV or CNN) and a working pipeline that accurately calculates and visualises metrics across Starts, Swim, and Turns.

Key challenges and difficulties anticipated include:

- **Data Annotation**: Obtaining sufficient, high-quality, frame-by-frame annotated footage for training the CNN-based lanerope model and for ground-truthing the pose/phase segmentation.
- **Generalisation**: Ensuring the lanerope calibration method is robust against varying pool environments, lighting conditions, and camera angles (beyond a fixed setup).
- **Symmetry Complexity**: Developing a sufficiently sensitive metric to accurately quantify left–right symmetry, given the inherent small variances in human movement.

6.1.1 Aims and Objectives

The primary aim of this project is to develop a Proof of Concept system that enables the extraction of swimming metrics from video footage. The specific objectives are:

1. To design, prototype, and create a modular, under-water marker design that maximizes underwater visibility.
2. To create and annotate a small dataset of swimming footage using the designed markers in the background.
3. To establish a pipeline for dynamic pixel-to-meter scale extraction using the markers for calibration.
4. To calculate simple metrics such as swimmer velocity and stroke length using the calibrated scale.

References

[1] W. Hsieh et al., *A Comprehensive Guide to Explainable AI: From Classical Models to LLMs*, 2024. DOI: 10.48550/ARXIV.2412.00800 Accessed: 30th Dec. 2025. [Online]. Available: https://arxiv.org/abs/2412.00800

[2] S. Ahmed, "Navigational system for visually impaired people in a swimming pool," Online, M.S. thesis, KTH Royal Institute of Technology, Stockholm, Sweden, 2019. [Online]. Available: https://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-269431

[3] N. Giulietti, A. Caputo, P. Chiariotti and P. Castellini, "SwimmerNET: Underwater 2D Swimmer Pose Estimation Exploiting Fully Convolutional Neural Networks," en, *Sensors*, vol. 23, no. 4, p. 2364, Jan. 2023, ISSN: 1424-8220. DOI: 10.3390/s23042364 Accessed: 5th Aug. 2025. [Online]. Available: https://www.mdpi.com/1424-8220/23/4/2364

[4] G. Jocher, J. Qiu and A. Chaurasia, *Ultralytics YOLO*, original-date: 2022-09-11T16:39:45Z, Jan. 2023. Accessed: 27th Apr. 2026. [Online]. Available: https://github.com/ultralytics/ultralytics

[5] C.-Y. Wang, A. Bochkovskiy and H.-Y. M. Liao, "Yolov7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2023, pp. 7464–7475.

[6] A. Kirillov et al., "Segment anything," in *Proceedings of the IEEE/CVF international conference on computer vision*, 2023, pp. 4015–4026.

[7] N. Ravi et al., "Sam 2: Segment anything in images and videos," *arXiv preprint arXiv:2408.00714*, 2024.

[8] Y. Xu, J. Zhang, Q. Zhang and D. Tao, "Vitpose: Simple vision transformer baselines for human pose estimation," *Advances in neural information processing systems*, vol. 35, pp. 38 571–38 584, 2022.

[9] A. Dosovitskiy et al., "An image is worth 16x16 words: Transformers for image recognition at scale," *arXiv preprint arXiv:2010.11929*, 2020.

[10] R. Hartley and A. Zisserman, *Multiple view geometry in computer vision*. Cambridge university press, 2003.

[11] R. E. Kalman, "A new approach to linear filtering and prediction problems," 1960.

[12] D. Zhou, J. W. L. Keogh, Y. Ma, R. K. Y. Tong, A. R. Khan and N. R. Jennings, "Artificial intelligence in sport: A narrative review of applications, challenges and future trends," *J. Sports Sci.*, 2025.

[13] A. Agrawal, S. Ramalingam, Y. Taguchi and V. Chari, "A theory of multi-layer flat refractive geometry," in *2012 IEEE conference on computer vision and pattern recognition*, IEEE, 2012, pp. 3346–3353.

[14] C. Holzmann and M. Hochgatterer, "Measuring distance with mobile phones using single-camera stereo vision," in *2012 32nd International conference on distributed computing systems workshops*, IEEE, 2012, pp. 88–93. Accessed: 5th Nov. 2025. [Online]. Available: https://ieeexplore.ieee.org/abstract/document/6258139/

[15] A. Al-Majnoni, J. Al-Sahli, D. Al-Ahmady, A. Al-Mutairi, A. Alsini and M. Alharbi, "Moar: A Swimmer Motion Swimming Style Identification Model using Deep Learning," en, *Engineering, Technology & Applied Science Research*, vol. 15, no. 1, pp. 19 295–19 302, Feb. 2025, ISSN: 1792-8036. DOI: 10.48084/etasr.9309 Accessed: 5th Aug. 2025. [Online]. Available: https://etasr.com/index.php/ETASR/article/view/9309

[16] T. Tran et al., "Analyzing swimming performance using drone captured aerial videos," in *Proceedings of the Micro Aerial Vehicle Networks, Systems, and Applications (DroNet '24)*, Tokyo, Japan: ACM, 2024. DOI: 10.1145/3663464

[17] J. W. Tukey et al., *Exploratory data analysis*. Springer, 1977, vol. 2.

[18] C. Leys, C. Ley, O. Klein, P. Bernard and L. Licata, "Detecting outliers: Do not use standard deviation around the mean, use absolute deviation around the median," *Journal of experimental social psychology*, vol. 49, no. 4, pp. 764–766, 2013.

Appendix A Sample A

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Mauris sed ipsum risus. Nulla aliquet quis quam sed eleifend. Donec rutrum, dolor id vulputate pharetra, nulla tortor laoreet nisl, pellentesque dapibus velit dolor suscipit purus. Phasellus vitae eleifend sem. Integer ultricies ex in neque pellentesque, vitae facilisis orci aliquam. In pellentesque mollis turpis, eu tristique lacus eleifend nec. Vestibulum orci neque, rhoncus vitae convallis eu, suscipit quis dui. Nulla libero elit, porta sit amet sagittis vel, placerat sit amet tortor. Aliquam hendrerit dolor sit amet sollicitudin ornare. Aliquam placerat sodales est, in vestibulum nisl efficitur in. Nulla venenatis aliquam sem, at volutpat nisl pellentesque eleifend. Praesent vitae euismod nulla, eget vehicula turpis. Duis quis tellus vitae nisi tempus tincidunt.

Nam quis aliquet nisi, non pharetra ligula. Phasellus pulvinar mattis neque, nec interdum justo condimentum hendrerit. Mauris fermentum venenatis faucibus. Pellentesque egestas eleifend libero, quis placerat ante fermentum et. Suspendisse accumsan gravida rhoncus. Vestibulum auctor sodales vehicula. Pellentesque a urna et elit placerat laoreet a quis turpis. Morbi ut sem at nunc posuere malesuada vel nec libero. Nam et urna suscipit, bibendum diam sed, aliquet leo. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut pretium mauris et nulla malesuada, mattis tristique lectus molestie. Aenean accumsan iaculis quam, eget varius libero placerat eu. Fusce mauris justo, vulputate a sollicitudin a, malesuada at est. Sed ac augue elit. Maecenas massa lorem, tincidunt vitae neque et, maximus dictum tellus.

Appendix B Sample B

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Mauris sed ipsum risus. Nulla aliquet quis quam sed eleifend. Donec rutrum, dolor id vulputate pharetra, nulla tortor laoreet nisl, pellentesque dapibus velit dolor suscipit purus. Phasellus vitae eleifend sem. Integer ultricies ex in neque pellentesque, vitae facilisis orci aliquam. In pellentesque mollis turpis, eu tristique lacus eleifend nec. Vestibulum orci neque, rhoncus vitae convallis eu, suscipit quis dui. Nulla libero elit, porta sit amet sagittis vel, placerat sit amet tortor. Aliquam hendrerit dolor sit amet sollicitudin ornare. Aliquam placerat sodales est, in vestibulum nisl efficitur in. Nulla venenatis aliquam sem, at volutpat nisl pellentesque eleifend. Praesent vitae euismod nulla, eget vehicula turpis. Duis quis tellus vitae nisi tempus tincidunt.

Nam quis aliquet nisi, non pharetra ligula. Phasellus pulvinar mattis neque, nec interdum justo condimentum hendrerit. Mauris fermentum venenatis faucibus. Pellentesque egestas eleifend libero, quis placerat ante fermentum et. Suspendisse accumsan gravida rhoncus. Vestibulum auctor sodales vehicula. Pellentesque a urna et elit placerat laoreet a quis turpis. Morbi ut sem at nunc posuere malesuada vel nec libero. Nam et urna suscipit, bibendum diam sed, aliquet leo. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut pretium mauris et nulla malesuada, mattis tristique lectus molestie. Aenean accumsan iaculis quam, eget varius libero placerat eu. Fusce mauris justo, vulputate a sollicitudin a, malesuada at est. Sed ac augue elit. Maecenas massa lorem, tincidunt vitae neque et, maximus dictum tellus.